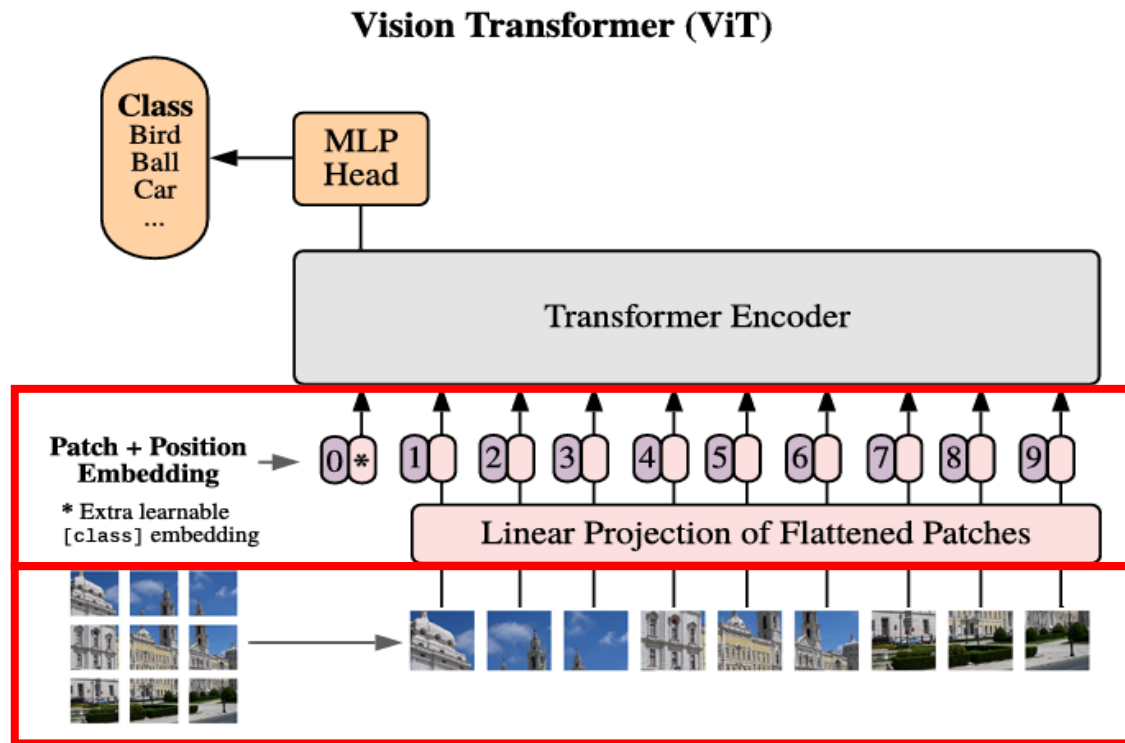


# ViT Short Tutorial

James

# Splitting images into patches



Source: "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale"

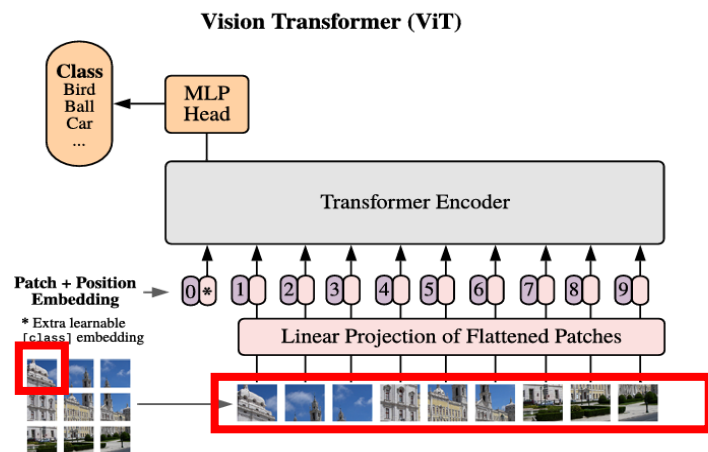
<https://arxiv.org/abs/2010.11929>

Q1: why do we divide it into small patches?

Main reason is that **image is too big for Transformers.**

Typical Transformers operates on a sequence of **tokens of a maximum length 512**. But a small image in CiFar10 has shape  $32 * 32 * 3$ , which after flattening gives 3072 length of tokens.

# How to Split?



$in\_channel = 3$   
 $img\_size = 224$   
 $patch\_size = 16$   
 $batch\_size = 512$

$$embed\_dim = (patch\_size^2) \times in\_channel = 16^2 \times 3 = 768$$

$nn.Conv2d($   
 $in\_channel = in\_channel,$   
 $out\_channel = embed\_dim,$   
 $kernel\_size = patch\_size,$   
 $stride = patch\_size)$

① (512, 3, 224, 224)

$$\begin{aligned}
 \text{output size } H' &= \left\lfloor \frac{H + 2 \times \text{padding} - \text{Kernel}}{\text{stride}} \right\rfloor + 1 \\
 &= \left\lfloor \frac{224 - 16}{16} \right\rfloor + 1 \\
 &= 14
 \end{aligned}$$

② (512, 768, 14, 14)

$nn.Flatten(2)$

$$\begin{aligned}
 num\_patches &= H' \times H' \\
 &= 14 \times 14 = 196
 \end{aligned}$$

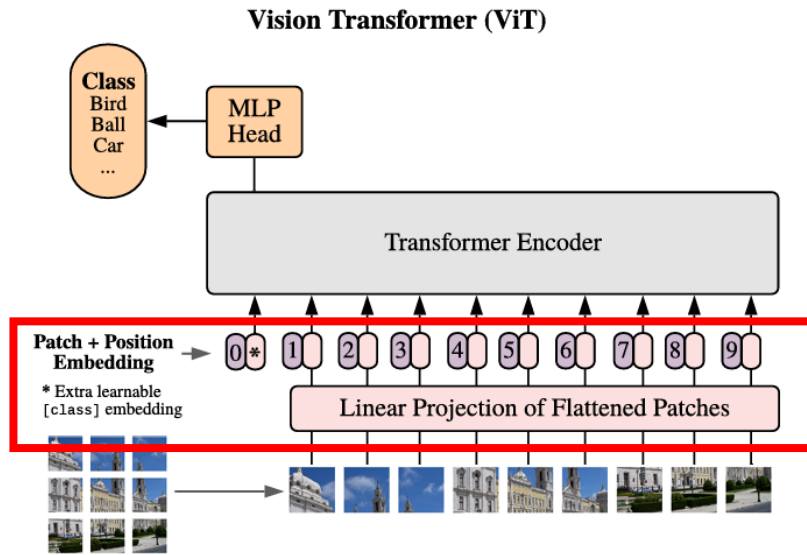
③ (512, 768, 196)

$nn.permute(0, 3, 1, 2)$

④ (512, 196, 768)

From a (3, 224, 224) image  
 divide into 196 patches  
 of size 768 dimension

# ViT-step2 cls\_token and position embedding

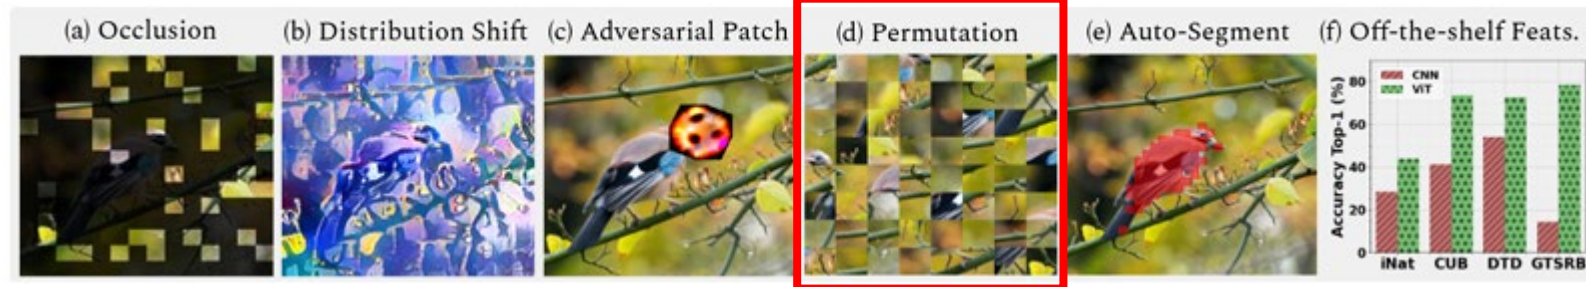


Operation	Size
Input_size	(512, 3, 224, 224)
nn.Conv2d	(512, 768, 14, 14)
nn.Flatten(2)	(512, 768, 196)
nn.permute(0, 2, 1)	(512, 196, 768)
Torch.cat([cls_token, x], dim=1)	(512, 197, 768)
$X = x + \text{self.position\_embedding}$	(512, 197, 768)
Transformer block	(512, 197, 768)
MLP-Head	(512, 10)

```
def forward(self, x):  
    # initiation cls token, and match the batch size of x in the 0th dimension  
    cls_token = self.cls_token.expand(x.shape[0], -1, -1)  
    # after flatten we need to permute the dimension  
    x = self.patcher(x).permute(0, 2, 1)  
    # concat cls token  
    x = torch.cat([x, cls_token], dim=1)  
    # add position embedding  
    x = x + self.position_embedding  
    x = self.dropout(x)  
    return x
```


# Transformer Encoder v.s. CNN

## Position Embedding



Source: "Intriguing Properties of Vision Transformers" <https://arxiv.org/abs/2105.10497>

```
self.position_embedding = nn.Parameter(torch.randn(1, num_patches+1, embed_dim), requires_grad=True)
```



Randomness Position Embedding **works**  
Because **position\_embedding assumes no prior information**. In other words, each patch position is learnt from training.

Thanks for watching